Interpretable Machine Learning

Counterfactual Explanations

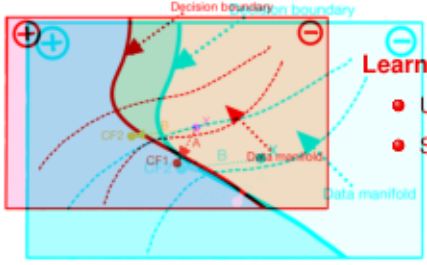

Learning goals

- Understand the motivation behind CEs
- See the mathematical foundation of CEs

EXAMPLE: CREDIT RISK APPLICATION

- $\mathbf{x}$: customer and credit information
- y: grant or reject credit

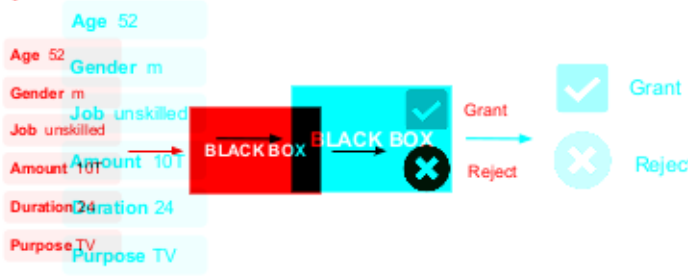

Questions:

- Why was the credit rejected?
- Is it a fair decision?
- **How should x be changed so that the credit is accepted?**

EXAMPLE: CREDIT RISK APPLICATION

Counterfactual Explanations provide answers in the form of "What-If"-scenarios.

Counterfactual Explanations provide answers in the form of "What-If"-scenarios.

"If the person was more skilled and the credit amount had been reduced to $8.000, the credit would have been granted."

"If the person was more skilled and the credit amount had been reduced to $8.000, the credit would have been granted."

Interpretable Machine Learning – 2 / 13

COUNTERFACTUAL EXPLANATIONS: MAIN IDEA

- Counterfactual explanations == counterfactuals == CEs

Interpretable Machine Learning – 3 / 13

COUNTERFACTUAL EXPLANATIONS: MAIN IDEA

- Counterfactual explanations == counterfactuals == CEs
- Explain particular predictions of an ML model by presenting an alternative input whose prediction equals a desired outcome

Interpretable Machine Learning – 3 / 13

COUNTERFACTUAL EXPLANATIONS: MAIN IDEA

- Counterfactual explanations == counterfactuals == CEs
- Explain particular predictions of an ML model by presenting an alternative input whose prediction equals a desired outcome
- Represent **close neighbors** of a data point we are interested in, but belonging to the **desired outcome**

Interpretable Machine Learning – 3 / 13

COUNTERFACTUAL EXPLANATIONS: MAIN IDEA

- Counterfactual explanations == counterfactuals == CEs
- Explain particular predictions of an ML model by presenting an alternative input whose prediction equals a desired outcome
- Represent **close neighbors** of a data point we are interested in, but belonging to the **desired outcome**
- Reveal which minimal changes to the input are sufficient to receive a different outcome
  ⇝ Useful if there is a chance to change the input features (e.g., by changing behaviour)

Counterfactual explanations == counterfactuals == CEs
Explain particular predictions of an ML model by presenting an alternative input whose prediction equals a desired outcome
Represent **close neighbors** of a data point we are interested in, but belonging to the **desired outcome**
Reveal which minimal changes to the input are sufficient to receive a different outcome
⇝ Useful if there is a chance to change the input features (e.g., by changing behaviour)

Interpretable Machine Learning – 3 / 13

COUNTERFACTUAL EXPLANATIONS: MAIN IDEA

- Counterfactual explanations == counterfactuals == CEs
- Explain particular predictions of an ML model by presenting an alternative input whose prediction equals a desired outcome
- Represent **close neighbors** of a data point we are interested in, but belonging to the **desired outcome**
- Reveal which minimal changes to the input are sufficient to receive a different outcome
  ⇝ Useful if there is a chance to change the input features (e.g., by changing behaviour)
- The targeted audience of CEs are often end-users

AIMS & ROLES

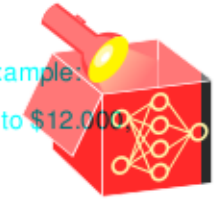

CEs can serve various purposes, the user can decide what to learn from them. For example:

"If the person had been **one year older** and the **credit amount had been increased** to $12.000,
the credit would have been granted."

Interpretable Machine Learning – 4 / 13

AIMS & ROLES

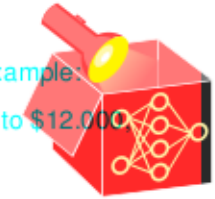

CEs can serve various purposes, the user can decide what to learn from them. For example:

"If the person had been **one year older** and the **credit amount had been increased** to $12.000, the credit would have been granted."

- **Guidance for future actions:**

  Ok, I will apply again next year for the higher amount.

CEs can serve various purposes, the user can decide what to learn from them. For example:

"If the person had been **one year older** and the **credit amount had been increased** to $12.000, the credit would have been granted."

- **Guidance for future actions:**

  Ok, I will apply again next year for the higher amount.

Interpretable Machine Learning – 4 / 13

AIMS & ROLES

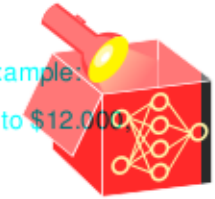

CEs can serve various purposes, the user can decide what to learn from them. For example:

"If the person had been **one year older** and the **credit amount had been increased** to $12.000,
the credit would have been granted."

- **Guidance for future actions:**
  Ok, I will apply again next year for the higher amount.
- **Provide reasons:**
  Interesting, I did not know that age plays a role in loan applications.

Interpretable Machine Learning – 4 / 13

AIMS & ROLES

CEs can serve various purposes, the user can decide what to learn from them. For example:

"If the person had been **one year older** and the **credit amount had been increased** to $12.000,
the credit would have been granted."

- **Guidance for future actions:**
  Ok, I will apply again next year for the higher amount.
- **Provide reasons:**
  Interesting, I did not know that age plays a role in loan applications.
- **Provide grounds to contest the decision:**
  How dare you, I do not want to be discriminated for my age in an application.

CEs can serve various purposes, the user can decide what to learn from them. For example:
"If the person had been **one year older** and the **credit amount had been increased** to $12.000,
the credit would have been granted."
- **Guidance for future actions:**
  Ok, I will apply again next year for the higher amount.
- **Provide reasons:**
  Interesting, I did not know that age plays a role in loan applications.
- **Provide grounds to contest the decision:**
  How dare you, I do not want to be discriminated for my age in an application.

Interpretable Machine Learning – 4 / 13

AIMS & ROLES

CEs can serve various purposes, the user can decide what to learn from them. For example:

"If the person had been **one year older** and the **credit amount had been increased** to $12.000, the credit would have been granted."

- **Guidance for future actions:**
  Ok, I will apply again next year for the higher amount.
- **Provide reasons:**
  Interesting, I did not know that age plays a role in loan applications.
- **Provide grounds to contest the decision:**
  How dare you, I do not want to be discriminated for my age in an application.
- **Detect model biases:**
  There is a bug, an increase in amount should not increase approval rates.

PHILOSOPHICAL BASIS

Counterfactuals have a long-standing tradition in analytic philosophy

⇝ According to [Lewis(1973)], a **counterfactual conditional** is a statement of the form:

"If S was the case, Q would have been the case."

"If S was the case, Q would have been the case."

Interpretable Machine Learning – 5 / 13

PHILOSOPHICAL BASIS

Counterfactuals have a long-standing tradition in analytic philosophy

⇝ According to [Lewis(1973)], a **counterfactual conditional** is a statement of the form:

"If S was the case, Q would have been the case."

- S is an event that must relate to a past event that didn't occur
- S is an event that must relate to a past event that didn't occur

  ⇝ counterfactuals run **contrary** to the **facts**

Interpretable Machine Learning – 5 / 13

PHILOSOPHICAL BASIS

Counterfactuals have a long-standing tradition in analytic philosophy

According to Lewis (1973), a **counterfactual conditional** is a statement of the form:

"If S was the case, Q would have been the case."

- S is an event that must relate to a past event that didn't occur
- counterfactuals run **contrary** to the **facts**
- Above statement is true, if in all possible worlds most similar to the actual world where S had been the case, Q would have been the case

Interpretable Machine Learning – 5 / 13

PHILOSOPHICAL BASIS

Counterfactuals have a long-standing tradition in analytic philosophy

⇝ According to Lewis (1973), a **counterfactual conditional** is a statement of the form:

"If S was the case, Q would have been the case."

- S is an event that must relate to a past event that didn't occur
- ⇝ counterfactuals run **contrary** to the **facts**
- Above statement is true, if in all possible worlds most similar to the actual world where S had been the case, Q would have been the case
- A world is similar to another if laws are maximally preserved between the worlds and only a few facts are changed

Interpretable Machine Learning – 5 / 13

PHILOSOPHICAL BASIS

- Counterfactuals have largely been studied to explain causal dependence

Interpretable Machine Learning – 6 / 13

PHILOSOPHICAL BASIS

- Counterfactuals have largely been studied to explain causal dependence
- Causal dependence underlies the explanatory power
  - ⇝ good CEs point to critical causal factors that drove the algorithmic decision

Interpretable Machine Learning – 6 / 13

PHILOSOPHICAL BASIS

- Counterfactuals have largely been studied to explain causal dependence
- Causal dependence underlies the explanatory power
  ⇝ good CEs point to critical causal factors that drove the algorithmic decision
- If maximal closeness is relaxed, causally irrelevant factors can become part of the explanation
  ⇝ e.g., decreasing loan amount by $20.000 and being one year older is recommended by the explainer although only loan amount might be causally relevant

Interpretable Machine Learning – 6 / 13

PHILOSOPHICAL BASIS

- Counterfactuals have largely been studied to explain causal dependence
- Causal dependence underlies the explanatory power
  ⇝ good CEs point to critical causal factors that drove the algorithmic decision
- If maximal closeness is relaxed, causally irrelevant factors can become part of the explanation
  ⇝ e.g., decreasing loan amount by $20.000 and being one year older is recommended by the explainer although only loan amount might be causally relevant
- CEs are often contrastive, i.e., they explain a decision by referring to an alternative outcome
  ⇝ e.g., if the loan applicant was 30 instead of 60 years old, the approved loan would have been over $100.000 instead of $40.000

Interpretable Machine Learning – 6 / 13

MATHEMATICAL PERSPECTIVE

Terminology:

- $\mathbf{x}$: original/factual datapoint whose prediction we want to explain
- y': desired prediction ($y' \neq 1000$ or $y' =$ "grant credit") or interval ($y' = [1000, \infty[$)

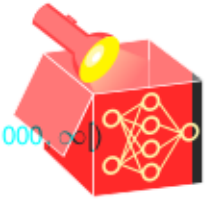

Interpretable Machine Learning – 7 / 13

MATHEMATICAL PERSPECTIVE

Terminology:

- $\mathbf{x}$: original/factual datapoint whose prediction we want to explain
- y': desired prediction ($y' = 1000$ or $y' =$ "grant credit") or interval ($y' = [1000, \infty[$)

A **valid** counterfactual $\mathbf{x}'$ is a datapoint:

1. whose prediction $\hat{f}(\mathbf{x}')$ is equal to the desired prediction y'
2. that is maximally close to the original datapoint $\mathbf{x}$

Interpretable Machine Learning – 7 / 13

MATHEMATICAL PERSPECTIVE

Terminology:

- $\mathbf{x}$: original/factual datapoint whose prediction we want to explain
- $y' \in \mathbb{R}$: desired prediction ($y' = 1000$ or $y' =$ "grant credit") or interval ($y' = [1000, \infty[$)

A **valid** counterfactual $\mathbf{x}'$ is a datapoint:

1. whose prediction $\hat{f}(\mathbf{x}')$ is equal to the desired prediction y'
2. that is maximally close to the original datapoint $\mathbf{x}$

Reformulate these two objectives (denoted by o_1 and o_2) as optimization problem:

$$\arg\min_{\mathbf{x}'} \lambda_1 o_p(\hat{f}(\mathbf{x}'), y') + \lambda_2 o_f(\mathbf{x}', \mathbf{x})$$

- λ_1 and λ_2 balance the two objectives
- Choice of o_p (distance on prediction space) and of o_f (distance on feature space) is crucial

MATHEMATICAL PERSPECTIVE

Dandl et al. (2020)

- Regression: o_p could be the L_1-distance $o_p(\hat{f}(\mathbf{x}'), y') = |\hat{f}(\mathbf{x}') - y'|$
- Classification: L_1-distance for scores and 0-1 Loss for labels, e.g., $o_p(\hat{f}(\mathbf{x}'), y') = \mathcal{I}_{\{\hat{f}(\mathbf{x}') \neq y'\}}$

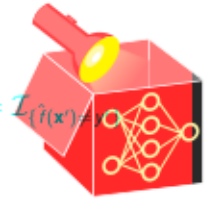

Interpretable Machine Learning – 8 / 13

MATHEMATICAL PERSPECTIVE

Dandl et al. (2020)

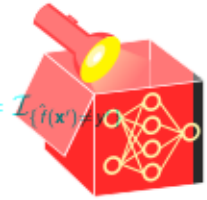

- Regression: o_p could be the L_1-distance $o_p(\hat{f}(\mathbf{x}'), y') = |\hat{f}(\mathbf{x}') - y'|$
- Classification: L_1-distance for scores and 0-1 Loss for labels, e.g., $o_p(\hat{f}(\mathbf{x}'), y') = \mathcal{I}_{\{\hat{f}(\mathbf{x}') \neq y'\}}$
- o_f could be the Gower distance (suitable for mixed feature space):

$$o_f(\mathbf{x}', \mathbf{x}) = d_G(\mathbf{x}', \mathbf{x}) = \frac{1}{p} \sum_{j=1}^{p} \delta_G(x'_j, x_j) \in [0, 1]$$

The value of δ_G depends on the feature type (numerical or categorical):

$$\delta_G(x'_j, x_j) = \begin{cases} \frac{1}{\widehat{R}_j} |x'_j - x_j| & \text{if } x_j \text{ is numerical} \\ \mathcal{I}_{\{x'_j \neq x_j\}} & \text{if } x_j \text{ is categorical} \end{cases}$$

with $\widehat{R}_j$ as the value range of feature j in the training dataset (to ensure that $\delta_G(x'_j, x_j) \in [0, 1]$)

Interpretable Machine Learning – 8 / 13

FURTHER OBJECTIVES

Additional constraints can improve the explanation quality of the corresponding CEs
⇝ popular constraints include sparsity and plausibility

Sparsity:

- End-users often prefer short over long explanations
  ⇝ counterfactuals should be **sparse**

Interpretable Machine Learning – 9 / 13

FURTHER OBJECTIVES

Additional constraints can improve the explanation quality of the corresponding CEs
⇝ popular constraints include sparsity and plausibility

Sparsity:

- End-users often prefer short over long explanations
  ⇝ counterfactuals should be **sparse**
- Objective o_1 can take the number of changed features into account (but does not have to)
  ⇝ e.g., the L_0- and the L_1-norm (similar to LASSO) can do this

Interpretable Machine Learning – 9 / 13

FURTHER OBJECTIVES

Additional constraints can improve the explanation quality of the corresponding CEs
⇝ popular constraints include sparsity and plausibility

Sparsity:

- End-users often prefer short over long explanations
  ⇝ counterfactuals should be **sparse**
- Objective o_c can take the number of changed features into account (but does not have to)
  ⇝ e.g., the L_0- and the L_1-norm (similar to LASSO) can do this
- Independently from o_c, sparsity in the changes can be additionally considered by another objective that counts the number of changed features via the L0-norm:

$$o_s(\mathbf{x}', \mathbf{x}) = \sum_{j=1}^{p} \mathcal{I}_{\{x_j' \neq x_j\}}$$

Interpretable Machine Learning – 9 / 13

FURTHER OBJECTIVES

Plausibility:

- CEs should suggest plausible alternatives
  - ⇝ e.g. not plausible to suggest to raise your income and get unemployed at the same time

Interpretable Machine Learning – 10 / 13

FURTHER OBJECTIVES

Plausibility:

- CEs should suggest plausible alternatives
  - ⇝ e.g., not plausible to suggest to raise your income and get unemployed at the same time
- CEs should be realistic and adhere to data manifold or originate from distribution of $\mathcal{X}$
  - ⇝ avoid unrealistic combinations of feature values

Interpretable Machine Learning – 10 / 13

FURTHER OBJECTIVES

Plausibility:

- CEs should suggest plausible alternatives
  - ⇝ e.g., not plausible to suggest to raise your income and get unemployed at the same time
- CEs should be realistic and adhere to data manifold or originate from distribution of $\mathcal{X}$
- Estimating joint distribution of training data is complex, especially for mixed feature spaces
  - ⇝ avoid unrealistic combinations of feature values
  - ⇝ Proxy: ensure that $\mathbf{x}'$ is close to training data $\mathbf{X}$

Interpretable Machine Learning – 10 / 13

FURTHER OBJECTIVES

Plausibility:

- CEs should suggest plausible alternatives
  - ⇝ e.g., not plausible to suggest to raise your income and get unemployed at the same time
- CEs should be realistic and adhere to data manifold or originate from distribution of $\mathcal{X}$
  - ⇝ avoid unrealistic combinations of feature values
- Estimating joint distribution of training data is complex, especially for mixed feature spaces
  - ⇝ Proxy: ensure that $\mathbf{x}'$ is close to training data $\mathbf{X}$

Example from Verma et al. (2020)

- Two possible paths for $\mathbf{x}$, originally classified to ⊖
- Two valid CEs in class ⊕: CF1 and CF2
- Path **A** for CF1 is shorter
- Path **B** for CF2 is longer but adheres to data manifold

FURTHER OBJECTIVES

To ensure plausibility, o_4 could, e.g., be the Gower distance of $\mathbf{x}'$ to its nearest data point of the training dataset which we denote $\mathbf{x}^{[1]}$:

$$o_4(\mathbf{x}', \mathbf{X}) = d_G(\mathbf{x}', \mathbf{x}^{[1]}) = \frac{1}{p}\sum_{j=1}^{p} \delta_G(x_j', x_j^{[1]})$$

We can extend the previous optimization problem by adding o_s (for sparsity) and o_4 (for plausibility):

$$\arg\min_{\mathbf{x}'} \lambda_1 o_p(\hat{f}(\mathbf{x}'), y') + \lambda_2 o_f(\mathbf{x}', \mathbf{x}) + \lambda_3 o_s(\mathbf{x}', \mathbf{x}) + \lambda_4 o_4(\mathbf{x}', \mathbf{X})$$

Interpretable Machine Learning – 11 / 13

REMARKS: THE RASHOMON EFFECT

Issue (Rashomon effect):

- Solution to the optimization problem might not be unique
- Many equally close CE might exist that obtain the desired prediction

  $\Rightarrow$ Many different equally good explanations for the same decision exist

Interpretable Machine Learning – 12 / 13

REMARKS: THE RASHOMON EFFECT

Issue (Rashomon effect):

- Solution to the optimization problem might not be unique
- Many equally close CE might exist that obtain the desired prediction
  $\Rightarrow$ Many different equally good explanations for the same decision exist

Possible solutions:

- Present all CEs for a given x (but: time and human processing capacity is limited)
- Focus on one or few CEs (but: by which criterion should they be selected?)

Interpretable Machine Learning – 12 / 13

REMARKS: THE RASHOMON EFFECT

Issue (Rashomon effect):

- Solution to the optimization problem might not be unique
- Many equally close CE might exist that obtain the desired prediction
  $\Rightarrow$ Many different equally good explanations for the same decision exist

Possible solutions:

- Present all CEs for a given x (but: time and human processing capacity is limited)
- Focus on one or few CEs (but: by which criterion should they be selected?)

Note:

- As the model is generally non-linear, inconsistent and diverse CEs can arise
  e.g. suggesting either an increase or decrease in credit duration (confuses the explainee)
- How to deal with the Rashomon effect is considered an open problem in IML

REMARKS: MODEL OR REAL-WORLD

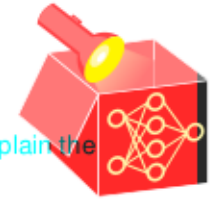

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  - $\rightsquigarrow$ Transfer of model explanations to explain real-world is generally not permitted

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  - $\rightsquigarrow$ Transfer of model explanations to explain real-world is generally not permitted

Interpretable Machine Learning – 13 / 13

REMARKS: MODEL OR REAL-WORLD

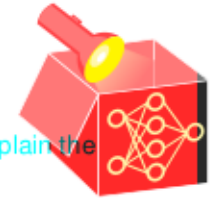

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  - $\rightsquigarrow$ Transfer of model explanations to explain real-world is generally not permitted
- Consider a CE that proposes to increase the feature age by 5 to obtain the loan
  - $\rightsquigarrow$ a loan applicant takes this information and applies 5 years later for the loan

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  - $\rightsquigarrow$ Transfer of model explanations to explain real-world is generally not permitted
- Consider a CE that proposes to increase the feature age by 5 to obtain the loan
  - $\rightsquigarrow$ a loan applicant takes this information and applies 5 years later for the loan

Interpretable Machine Learning – 13 / 13

REMARKS: MODEL OR REAL-WORLD

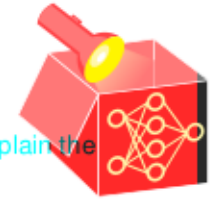

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  ⇝ Transfer of model explanations to explain real-world is generally not permitted
- Consider a CE that proposes to increase the feature age by 5 to obtain the loan
  ⇝ a loan applicant takes this information and applies 5 years later for the loan
- However, by then, many other feature values might have changed
  ⇝ not only age, also other causally dependent features e.g. job status might have changed
  ⇝ Karimi et al. (2020) avoid this by considering causal dependencies between features

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  ⇝ Transfer of model explanations to explain real-world is generally not permitted
- Consider a CE that proposes to increase the feature age by 5 to obtain the loan
  ⇝ a loan applicant takes this information and applies 5 years later for the loan
- However, by then, many other feature values might have changed
  ⇝ not only age, also other causally dependent features e.g. job status might have changed
  ⇝ Karimi et al. (2020) avoid this by considering causal dependencies between features

Interpretable Machine Learning – 13 / 13

REMARKS: MODEL OR REAL-WORLD

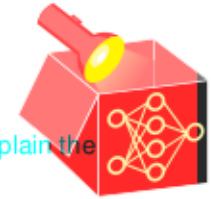

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  - ⇝ Transfer of model explanations to explain real-world is generally not permitted
- Consider a CE that proposes to increase the feature age by 5 to obtain the loan
  - ⇝ a loan applicant takes this information and applies 5 years later for the loan
- However, by then, many other feature values might have changed
  - ⇝ not only age, also other causally dependent features e.g. job status might have changed
  - ⇝ Karimi et al. (2020) avoid this by considering causal dependencies between features
- Also, the bank's algorithm might change and previous CEs are not applicable anymore

- Most CEs provide explanations of model predictions, but CEs might appear to explain the real-world for end-users
  - ⇝ Transfer of model explanations to explain real-world is generally not permitted
- Consider a CE that proposes to increase the feature age by 5 to obtain the loan
  - ⇝ a loan applicant takes this information and applies 5 years later for the loan
- However, by then, many other feature values might have changed
  - ⇝ not only age, also other causally dependent features e.g. job status might have changed
  - ⇝ Karimi et al. (2020) avoid this by considering causal dependencies between features
- Also, the bank's algorithm might change and previous CEs are not applicable anymore

Interpretable Machine Learning – 13 / 13