Interpretable Machine Learning

Rule-based Models

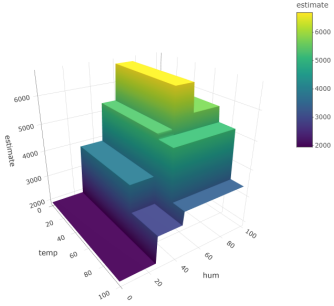

Learning goals

- Decision trees
- RuleFit
- Decision rules

DECISION TREES

▸ Breiman et al. (1984)

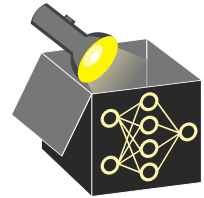

Idea: Partition data into axis-aligned regions via greedy search for feature cut points (minimizing a split criterion), then predict a constant mean c_m in each leaf region $\mathcal{R}_m$:

$$\hat{f}(x) = \sum_{m=1}^{M} c_m \mathbb{1}_{\{x \in \mathcal{R}_m\}}$$

DECISION TREES

▸ Breiman et al. (1984)

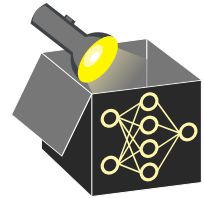

Idea: Partition data into axis-aligned regions via greedy search for feature cut points (minimizing a split criterion), then predict a constant mean c_m in each leaf region $\mathcal{R}_m$:

$$\hat{f}(x) = \sum_{m=1}^{M} c_m \mathbb{1}_{\{x \in \mathcal{R}_m\}}$$

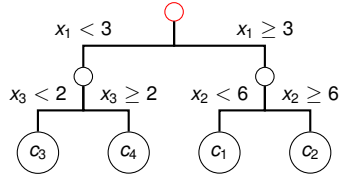

- Applicable to regression and classification
- Models interactions and non-linear effects
- Handles mixed feature spaces & missing values

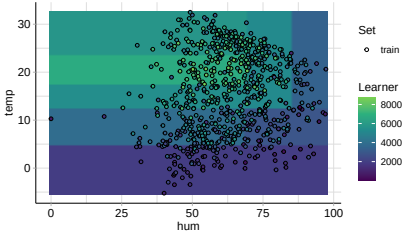

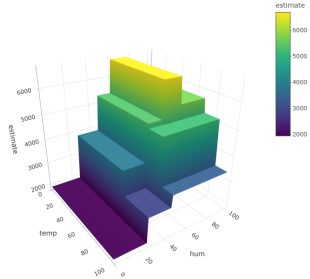

INTERPRETATION OF TREE-BASED MODELS

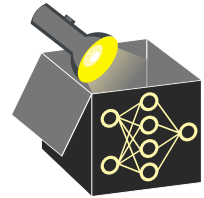

- Interpretation via path of decision rules along tree branches
- **Feature importance** (quantifies how often and how usefully x_j is used):
  - For each split on feature x_j, record the decrease in the split criterion
  - Aggregate this over the tree: sum or average over all splits involving x_j
  - Split criterion: variance (regression), Gini index / entropy (classification)

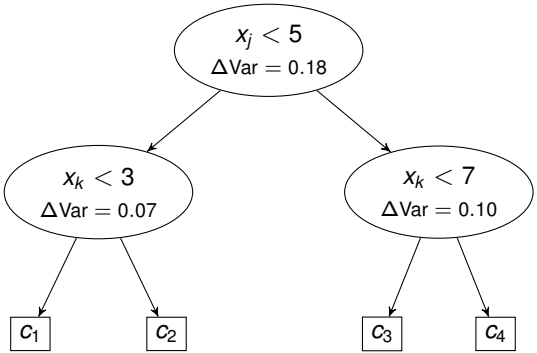

- Each ΔVar is assigned to the splitting feature
- Feature importance = sum of all ΔVar for that feature:

  x_j: 0.18

  x_k: $0.07 + 0.10 = 0.17$

DECISION TREES - EXAMPLE

- Fit decision tree with tree depth of 3 on bike data
- E.g., mean prediction for the first 105 days since 2011 is 1798
  $\rightsquigarrow$ Applies to $\hat{=}$15% of the data (leftmost branch)
- $\mathrm{days_since_2011}$: highest feature importance (explains most of variance)

Feature	Importance
days_since_2011	79.53
temp	17.55
hum	2.92

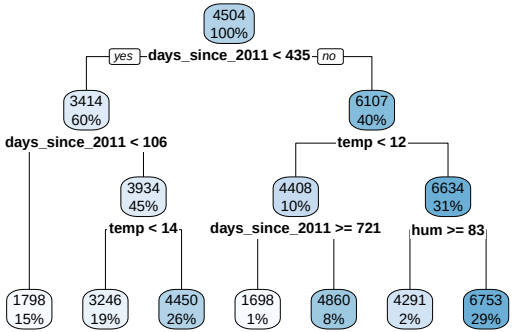

UNBIASED RECURSIVE PARTITIONING

▸ Hothorn et al. (2006) ▸ Zeileis et al. (2008) ▸ Strobl et al. (2007)

Problems with CART (Classification and Regression Trees):

1. Selection bias towards high-cardinal/continuous features
2. Splits on any improvement, regardless of significance $\rightsquigarrow$ prone to overfitting

UNBIASED RECURSIVE PARTITIONING

▸ Hothorn et al. (2006) ▸ Zeileis et al. (2008) ▸ Strobl et al. (2007)

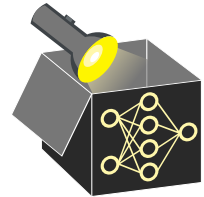

Problems with CART (Classification and Regression Trees):

1. Selection bias towards high-cardinal/continuous features
2. Splits on any improvement, regardless of significance $\rightsquigarrow$ prone to overfitting

Unbiased recursive partitioning via conditional inference trees (ctree) or model-based recursive partitioning (mob):

1. Separate selection of **feature used for splitting** and **split point**
2. Hypothesis test as stopping criteria

UNBIASED RECURSIVE PARTITIONING

▸ Hothorn et al. (2006) ▸ Zeileis et al. (2008) ▸ Strobl et al. (2007)

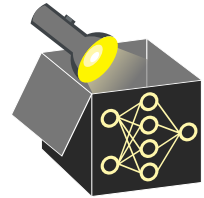

Problems with CART (Classification and Regression Trees):

1. Selection bias towards high-cardinal/continuous features
2. Splits on any improvement, regardless of significance $\rightsquigarrow$ prone to overfitting

Unbiased recursive partitioning via conditional inference trees (ctree) or model-based recursive partitioning (mob):

1. Separate selection of **feature used for splitting** and **split point**
2. Hypothesis test as stopping criteria

Example (selection bias):
Simulate data ($n = 200$) with $Y \sim N(0, 1)$ and 3 features of different cardinality independent from Y (repeat 500 times):

- $X_1 \sim Binom(n, \frac{1}{2})$
- $X_2 \sim M(n, (\frac{1}{4}, \frac{1}{4}, \frac{1}{4}, \frac{1}{4}))$
- $X_3 \sim M(n, (\frac{1}{8}, \frac{1}{8}, \frac{1}{8}, \frac{1}{8}, \frac{1}{8}, \frac{1}{8}, \frac{1}{8}, \frac{1}{8}))$

Which feature is selected in the first split?

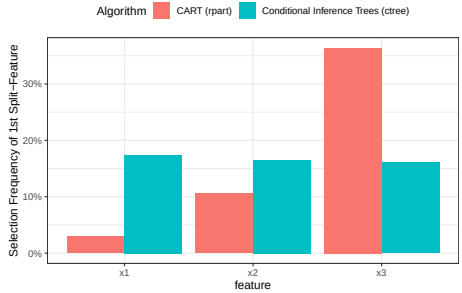

UNBIASED RECURSIVE PARTITIONING

Differences to CART:

- Two-step approach (1. find most significant split feature, 2. find best split point)
- Parametric model (e.g. LM instead of constant) can be fitted in leave nodes
- Significance of split (p-value) given in each node
- ctree and mob differ in hypothesis test used for selecting the split feature (independence test vs. fluctuation test) and how to find the best split point

UNBIASED RECURSIVE PARTITIONING

Differences to CART:

- Two-step approach (1. find most significant split feature, 2. find best split point)
- Parametric model (e.g. LM instead of constant) can be fitted in leave nodes
- Significance of split (p-value) given in each node
- ctree and mob differ in hypothesis test used for selecting the split feature (independence test vs. fluctuation test) and how to find the best split point

Example (ctree): Bike data (constant model in final nodes)

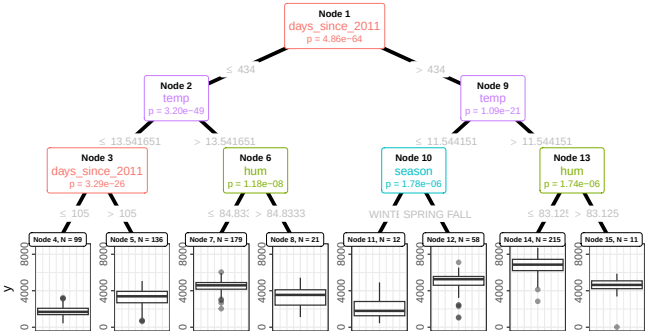

Train error (MSE):
758,844.0 (ctree)
742,244.4 (mob)

UNBIASED RECURSIVE PARTITIONING

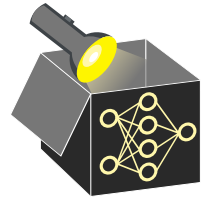

Differences to CART:

- Two-step approach (1. find most significant split feature, 2. find best split point)
- Parametric model (e.g. LM instead of constant) can be fitted in leave nodes
- Significance of split (p-value) given in each node
- ctree and mob differ in hypothesis test used for selecting the split feature (independence test vs. fluctuation test) and how to find the best split point

Example (mob): Bike data (linear model with temp in final nodes)

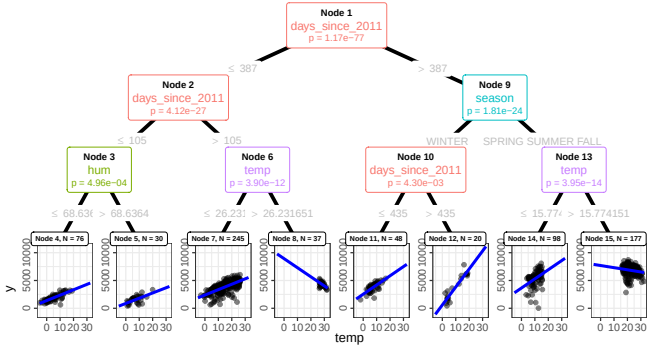

Train error (MSE):
758,844.0 (ctree)
742,244.4 (mob)

OTHER RULE-BASED MODELS

Decision Rules

Holte 1993

- Flat list of simple "if – then" statements
  $\rightsquigarrow$ very intuitive and easy-to-interpret
- Mainly devised for classification
  (support for regression is limited)
- Numeric features are typically discretised

IF $x_1 \leq 2.3$ AND $x_4 =$ "A"	THEN	y = 1
ELSE IF $x_2 > 5.0$	THEN	y = 2
ELSE		y = 3

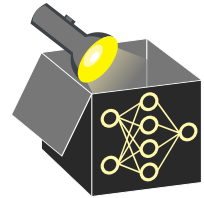

OTHER RULE-BASED MODELS

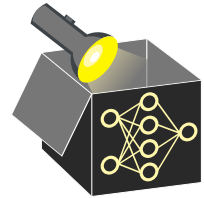

Decision Rules

Holte 1993

- Flat list of simple "if – then" statements
  $\rightsquigarrow$ very intuitive and easy-to-interpret
- Mainly devised for classification
  (support for regression is limited)
- Numeric features are typically discretised

IF $x_1 \leq 2.3$ AND $x_4 =$ "A" THEN y = 1
ELSE IF $x_2 > 5.0$ THEN y = 2
ELSE y = 3

RuleFit

Friedman & Popescu 2008

- Extract binary rules $r_m(\mathbf{x}) \in \{0, 1\}$ from many shallow trees (one per root-to-leaf path)
- Fit an L_1-regularized LM
  $\hat{f}(\mathbf{x}) = \beta_0 + \sum_m \beta_m r_m(\mathbf{x}) + \sum_j \gamma_j x_j$
- Regularization retains only a few rules
  $\Rightarrow$ sparse, non-linear, interaction-aware
- Coefficients relate to rule/feature importance

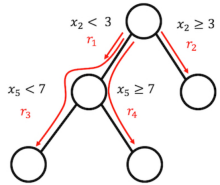

Molnar 2022