



Interpretable Machine Learning



Feature Effects

Accumulated Local Effect (ALE) plot



Learning goals

- Understand ALE plots
- Difference between ALE and PD plots


ACCUMULATED LOCAL EFFECTS (ALE)

► “Apley, Zhu” 2020

ALE plots estimate marginal effect of a feature by accumulating its local effects (integrating partial derivatives), evaluated in regions supported by the data.

Computation Steps:

- 1 **Estimate local effects** $\frac{\partial \hat{f}(x_S, \mathbf{x}_{-S})}{\partial x_S}$ (via finite differences)
⇒ Removes unwanted main effects of other features $\mathbf{x}_{-S}$ (unlike M plots)




ACCUMULATED LOCAL EFFECTS (ALE)

► “Apley, Zhu” 2020

ALE plots estimate marginal effect of a feature by accumulating its local effects (integrating partial derivatives), evaluated in regions supported by the data.

Computation Steps:

- 1 **Estimate local effects** $\frac{\partial \hat{f}(x_S, \mathbf{x}_{-S})}{\partial x_S}$ (via finite differences)
⇒ Removes unwanted main effects of other features $\mathbf{x}_{-S}$ (unlike M plots)
- 2 **Average local effects** over conditional distr. $\mathbb{P}(\mathbf{x}_{-S} | x_S)$ similar to M plots
⇒ Avoids extrapolation (unlike PD plots)




ACCUMULATED LOCAL EFFECTS (ALE)

► “Apley, Zhu” 2020

ALE plots estimate marginal effect of a feature by accumulating its local effects (integrating partial derivatives), evaluated in regions supported by the data.

Computation Steps:

- 1 **Estimate local effects** $\frac{\partial \hat{f}(x_S, \mathbf{x}_{-S})}{\partial x_S}$ (via finite differences)
⇒ Removes unwanted main effects of other features $\mathbf{x}_{-S}$ (unlike M plots)
- 2 **Average local effects** over conditional distr. $\mathbb{P}(\mathbf{x}_{-S} | x_S)$ similar to M plots
⇒ Avoids extrapolation (unlike PD plots)
- 3 **Accumulate:** Integrate averaged local effects up to a specific $x \in \mathcal{X}_S$
⇒ Reconstructs main effect of x_S




FIRST ORDER ALE FUNCTION

Uncentered ALE Function evaluated at $x \in \mathcal{X}_S$ (domain of feature x_S):

$$\tilde{f}_{S,\text{ALE}}(x) = \underbrace{\int_{z_0}^x}_{(3)} \underbrace{\mathbb{E}_{\mathbf{x}_{-S} | x_S = z_S}}_{(2) \text{ average locally}} \left(\underbrace{\frac{\partial \hat{f}(z_S, \mathbf{x}_{-S})}{\partial z_S}}_{(1) \text{ local effect}} \right) dz_S = \int_{z_0}^x \int \frac{\partial \hat{f}(z_S, \mathbf{x}_{-S})}{\partial z_S} d\mathbb{P}(\mathbf{x}_{-S} | z_S) dz_S$$


- x_S is feature of interest, with minimum value $z_0 = \min(x_S)$
- z_S is integration variable ranging over $\mathcal{X}_S$, used to evaluate local effects
- $\mathbf{x}_{-S}$ denotes all other features (complement of S)



FIRST ORDER ALE FUNCTION

Uncentered ALE Function evaluated at $x \in \mathcal{X}_S$ (domain of feature x_S):

$$\tilde{f}_{S,\text{ALE}}(x) = \underbrace{\int_{z_0}^x}_{(3)} \underbrace{\mathbb{E}_{\mathbf{x}_{-S} | x_S = z_S}}_{(2) \text{ average locally}} \left(\underbrace{\frac{\partial \hat{f}(z_S, \mathbf{x}_{-S})}{\partial z_S}}_{(1) \text{ local effect}} \right) dz_S = \int_{z_0}^x \int \frac{\partial \hat{f}(z_S, \mathbf{x}_{-S})}{\partial z_S} d\mathbb{P}(\mathbf{x}_{-S} | z_S) dz_S$$






- x_S is feature of interest, with minimum value $z_0 = \min(x_S)$
- z_S is integration variable ranging over $\mathcal{X}_S$, used to evaluate local effects
- $\mathbf{x}_{-S}$ denotes all other features (complement of S)

Centering (to ensure identifiability):

$$f_{S,\text{ALE}}(x) = \tilde{f}_{S,\text{ALE}}(x) - \underbrace{\int \tilde{f}_{S,\text{ALE}}(x_S) d\mathbb{P}(x_S)}_{\text{constant shift to mean zero}}$$

ALE ESTIMATION: ILLUSTRATION






- **Motivation:** Partial derivatives are not well-defined for all models (e.g., tree-based methods). $\Rightarrow$ Use finite differences within intervals instead.
- Partition the feature range of x_S into K intervals (vertical lines)
 - Define intervals:

$$x_S \in [\min(x_S), \max(x_S)] \Rightarrow x_S \in [z_0, z_{1,S}] \cup [z_{1,S}, z_{2,S}] \cup \dots \cup [z_{K-1,S}, z_{K,S}]$$




- *Equidistant*: preserves resolution
- *Quantile-based*: balances sample size per interval

ALE ESTIMATION: ILLUSTRATION



- For each observation in k -th interval, i.e., $\{i : x_S^{(i)} \in [z_{k-1,S}, z_{k,S}]\}$:
 - Replace $x_S^{(i)}$ with **upper/lower interval bounds**, keeping $\mathbf{x}_{-S}^{(i)}$ fixed
 - Compute obs.-wise finite difference of i -th obs. in k -th interval
 $\leadsto \hat{f}(z_{k,S}, \mathbf{x}_{-S}^{(i)}) - \hat{f}(z_{k-1,S}, \mathbf{x}_{-S}^{(i)})$ (approximates local effect)

ALE ESTIMATION: ILLUSTRATION




- For each observation in k -th interval, i.e., $\{i : x_S^{(i)} \in [z_{k-1,S}, z_{k,S}]\}$:
 - Replace $x_S^{(i)}$ with **upper/lower interval bounds**, keeping $\mathbf{x}_{-S}^{(i)}$ fixed
 - Compute obs.-wise finite difference of i -th obs. in k -th interval
 $\rightsquigarrow \hat{f}(z_{k,S}, \mathbf{x}_{-S}^{(i)}) - \hat{f}(z_{k-1,S}, \mathbf{x}_{-S}^{(i)})$ (approximates local effect)
- Average these finite differences over all observations in each interval
 $\rightsquigarrow$ Approximates **inner integral** $\mathbb{E}_{\mathbf{x}_{-S} | x_S = z_S} \left[\partial \hat{f} / \partial z_S \right]$
- Accumulate these averages from z_0 to the point of interest $x \in \mathcal{X}_S$
 $\rightsquigarrow$ Approximates **outer integral** over $z_S \in [z_0, x]$
 $\Rightarrow$ uncentered ALE function

ALE ESTIMATION: FORMULA

Estimated uncentered ALE: For a point $x \in \mathcal{X}_S$, define:

$$\hat{f}_{S,ALE}(x) = \sum_{k=1}^{k_S(x)} \frac{1}{n_S(k)} \sum_{i: x_S^{(i)} \in [z_{k-1,S}, z_{k,S}]} \left[\hat{f}(z_{k,S}, \mathbf{x}_{-S}^{(i)}) - \hat{f}(z_{k-1,S}, \mathbf{x}_{-S}^{(i)}) \right]$$

- $[z_{k-1,S}, z_{k,S}]$: k -th interval of feat. x_S with interval bounds $z_{k-1,S}$ and $z_{k,S}$
- $k_S(x)$: index of the interval in which x lies
- $n_S(k)$: number of observations in interval k
- $\mathbf{x}_{-S}^{(i)}$: all other features held fixed for i -th observation



ALE ESTIMATION: FORMULA

Estimated uncentered ALE: For a point $x \in \mathcal{X}_S$, define:

$$\hat{f}_{S,ALE}(x) = \sum_{k=1}^{k_S(x)} \frac{1}{n_S(k)} \sum_{i: x_S^{(i)} \in [z_{k-1,S}, z_{k,S}]} \left[\hat{f}(z_{k,S}, \mathbf{x}_{-S}^{(i)}) - \hat{f}(z_{k-1,S}, \mathbf{x}_{-S}^{(i)}) \right]$$

- $[z_{k-1,S}, z_{k,S}]$: k -th interval of feat. x_S with interval bounds $z_{k-1,S}$ and $z_{k,S}$
- $k_S(x)$: index of the interval in which x lies
- $n_S(k)$: number of observations in interval k
- $\mathbf{x}_{-S}^{(i)}$: all other features held fixed for i -th observation


Centering: Ensure identifiability by subtracting mean uncentered ALE (c):

$$\hat{f}_{S,ALE}(x) = \hat{\tilde{f}}_{S,ALE}(x) - c, \quad c = \frac{1}{n} \sum_{i=1}^n \hat{\tilde{f}}_{S,ALE}(x_S^{(i)}).$$


Efficient centering (used in implementations): Use weighted trapezoidal averaging of interval-wise boundary values (avoids redundant re-evaluation at all n points):

$$c = \sum_{k=1}^K \frac{1}{2} \cdot \left(\hat{\tilde{f}}_{S,ALE}(z_{k-1,S}) + \hat{\tilde{f}}_{S,ALE}(z_{k,S}) \right) \cdot \frac{n_S(k)}{n}$$


Plotting: Visualize pairs $\left(z_{k,S}, \hat{f}_{S,ALE}(z_{k,S}) \right)$ for all interval boundaries $z_{k,S}$.



BIKE SHARING DATASET: FIRST ORDER ALE





- **Visual comparison:** PD plot (top) vs. First-order ALE plot (bottom)
- **Shape:** Similar trends in both plots; y-axis scale differs due to centering
- **Interpretation:** ALE accounts for feature dependencies and avoids extrapolation into unsupported regions
 - ↪ PD reflects model behavior in entire feature space ("true to the model")
 - ↪ ALE focuses on effects in data-supported regions ("true to the data")



BIKE SHARING DATASET: SECOND ORDER ALE

Unlike bivariate PD plots, 2nd-order ALE plots only estimate pure interaction between two features (1st-order effects are not included).




PD VS. ALE

PD:

$$f_{S,PD}(x_S) = \mathbb{E}_{\mathbf{x}_{-S}} \left(\hat{f}(x_S, \mathbf{x}_{-S}) \right)$$

ALE:

$$f_{S,ALE}(x) = \int_{z_0}^x \mathbb{E}_{\mathbf{x}_{-S} | x_S = z_S} \left(\frac{\partial \hat{f}(z_S, \mathbf{x}_{-S})}{\partial z_S} \right) dz - \text{const}$$



- Recall: PD directly averages predictions over marginal distribution of $\mathbf{x}_{-S}$
- ALE is faster: $O(2 \cdot n)$ model calls vs. $O(n \cdot g)$ for PD with g grid points
- Difference 1: ALE averages
 - **prediction changes** (via partial derivatives, estimated by finite differences)
 - over **conditional distribution** $\mathbb{P}(\mathbf{x}_{-S} | x_S = z_S)$


PD VS. ALE

PD:

$$f_{S,PD}(x_S) = \mathbb{E}_{\mathbf{x}_{-S}} \left(\hat{f}(x_S, \mathbf{x}_{-S}) \right)$$

ALE:

$$f_{S,ALE}(x) = \int_{z_0}^x \mathbb{E}_{\mathbf{x}_{-S} | x_S = z_S} \left(\frac{\partial \hat{f}(z_S, \mathbf{x}_{-S})}{\partial z_S} \right) dz - \text{const}$$



- Recall: PD directly averages predictions over marginal distribution of $\mathbf{x}_{-S}$
- ALE is faster: $O(2 \cdot n)$ model calls vs. $O(n \cdot g)$ for PD with g grid points
- Difference 1: ALE averages the
 - prediction changes (via partial derivatives, estimated by finite differences)
 - over conditional distribution $\mathbb{P}(\mathbf{x}_{-S} | x_S = z_S)$
- Difference 2: ALE integrates these partial deriv. over $z_S \in [z_0, x] \subseteq \mathcal{X}_S$
 $\rightsquigarrow$ isolates effect of x_S and removes main effect of other dependent feat.


PD VS. ALE

PD:

$$f_{S,PD}(x_S) = \mathbb{E}_{\mathbf{x}_{-S}} \left(\hat{f}(x_S, \mathbf{x}_{-S}) \right)$$

ALE:

$$f_{S,ALE}(x) = \int_{z_0}^x \mathbb{E}_{\mathbf{x}_{-S}|x_S=z_S} \left(\frac{\partial \hat{f}(z_S, \mathbf{x}_{-S})}{\partial z_S} \right) dz - \int \tilde{f}_{S,ALE}(x_S) d\mathbb{P}(x_S)$$



- Recall: PD directly averages predictions over marginal distribution of $\mathbf{x}_{-S}$
- ALE is faster: $O(2 \cdot n)$ model calls vs. $O(n \cdot g)$ for PD with g grid points
- Difference 1: ALE averages the
 - prediction changes (via partial derivatives, estimated by finite differences)
 - over conditional distribution $\mathbb{P}(\mathbf{x}_{-S}|x_S = z_S)$
- Difference 2: ALE integrates these partial deriv. over $z_S \in [z_0, x] \subseteq \mathcal{X}_S$
 $\rightsquigarrow$ isolates effect of x_S and removes main effect of other dependent feat.
- Difference 3: ALE is **centered** so that $\mathbb{E}_{x_S} (f_{S,ALE}(x)) = 0$