

Interpretable Machine Learning

Feature Importance

Shapley Additive Global Importance (SAGE)





Figure: Bike Sharing Dataset

Learning goals

- How SAGE fairly distributes importance
- Definition of SAGE value function
- Difference SAGE value function and SAGE values
- Marginal and Conditional SAGE



CHALLENGE: FAIR ATTRIBUTION OF IMPORTANCE

Recap:

- Data: $x_1, \dots, x_4$ uniformly sampled from $[-1, 1]$
- DGP: $y := x_1 x_2 + x_3 + \epsilon_Y$ with $\epsilon_Y \sim N(0, 1)$
- Model: $\hat{f}(x) \approx x_1 x_2 + x_3$

Although x_3 alone contributes as much to the prediction as x_1 and x_2 jointly, all three are considered equally relevant by PFI.

Reason: PFI assesses importance given that all remaining features are preserved. If we first permute x_1 and then x_2 , permutation of x_2 would have no effect on the performance (and vice versa).



SAGE IDEA

► "Covert et al." 2020


SAGE: Use Shapley values to compute a fair attribution of importance (via model performance)

Idea:

- Feature importance attribution can be regarded as cooperative game
 $\rightsquigarrow$ features jointly contribute to achieve a certain model performance
- Players: features
- Payoff to be fairly distributed: model performance
- Surplus contribution of a feature depends on the coalition of features that are already accessible by the model

Note:


- Similar idea (called SFIMP) was proposed in ► "Casalicchio et al." 2018
- Definition based on model refits was proposed in context of feature selection in
 ► "Cohen et al." 2007



SAGE - VALUE FUNCTION

Removal Idea: To deprive information of the non-coalition features $-S$ from the model, marginalize the prediction function over feats $-S$ to be “dropped”.

$$\hat{f}_S(x_S) = \mathbb{E}[\hat{f}(x) | X_S = x_S]$$



SAGE - VALUE FUNCTION


Removal Idea: To deprive information of the non-coalition features — S from the model, marginalize the prediction function over feats — S to be “dropped”.

$$\hat{f}_S(x_S) = \mathbb{E}[\hat{f}(x)|X_S = x_S]$$

SAGE value function:

$$v_{\hat{f}}(S) = \mathcal{R}(\hat{f}_{\emptyset}) - \mathcal{R}(\hat{f}_S), \text{ where } \mathcal{R}(\hat{f}_S) = \mathbb{E}_{Y, X_S}[L(y, \hat{f}_S(x_S))]$$

- ↪ Quantify the predictive power of a coalition S in terms of reduction in risk
- ↪ Risk of predictor $\hat{f}_S(x_S)$ is compared to the risk of the mean prediction $\hat{f}_{\emptyset}$



SAGE - VALUE FUNCTION

Removal Idea: To deprive information of the non-coalition features $-S$ from the model, marginalize the prediction function over feats $-S$ to be “dropped”.

$$\hat{f}_S(x_S) = \mathbb{E}[\hat{f}(x)|X_S = x_S]$$

SAGE value function:


$$v_{\hat{f}}(S) = \mathcal{R}(\hat{f}_{\emptyset}) - \mathcal{R}(\hat{f}_S), \text{ where } \mathcal{R}(\hat{f}_S) = \mathbb{E}_{Y, X_S}[L(Y, \hat{f}_S(x_S))]$$

- ↪ Quantify the predictive power of a coalition S in terms of reduction in risk
- ↪ Risk of predictor $\hat{f}_S(x_S)$ is compared to the risk of the mean prediction $\hat{f}_{\emptyset}$

Surplus contribution of feature x_j over coalition x_S :

$$v_{\hat{f}}(S \cup \{j\}) - v_{\hat{f}}(S) = \mathcal{R}(\hat{f}_S) - \mathcal{R}(\hat{f}_{S \cup \{j\}})$$


- ↪ Quantifies the added value of feature j when it is added to coalition S



SAGE - MARGINAL AND COND. SAMPLING

When computing the marginalized prediction $\hat{f}_S(x_S)$, the “dropped” features can be sampled from

- the marginal distribution $\mathbb{P}(x_{-S}) \Rightarrow$ marginal SAGE
- the conditional distribution $\mathbb{P}(x_{-S}|x_S) \Rightarrow$ conditional SAGE




SAGE - MARGINAL AND COND. SAMPLING

When computing the marginalized prediction $\hat{f}_S(x_S)$, the “dropped” features can be sampled from

- the marginal distribution $\mathbb{P}(x_{-S}) \Rightarrow$ marginal SAGE
- the conditional distribution $\mathbb{P}(x_{-S}|x_S) \Rightarrow$ conditional SAGE

Interpretation marginal sampling: $v(S)$ quantifies the reliance of the model on features x_S

- features x_S not being causal for the prediction $\Rightarrow v(S) = 0$



SAGE - MARGINAL AND COND. SAMPLING

When computing the marginalized prediction $\hat{f}_S(x_S)$, the “dropped” features can be sampled from


- the marginal distribution $\mathbb{P}(x_{-S}) \Rightarrow$ marginal SAGE
- the conditional distribution $\mathbb{P}(x_{-S}|x_S) \Rightarrow$ conditional SAGE

Interpretation marginal sampling: $v(S)$ quantifies the reliance of the model on features x_S

- features x_S not being causal for the prediction $\Rightarrow v(S) = 0$

Interpretation conditional sampling: $v(S)$ quantifies whether variables x_S contain prediction-relevant information (e.g. $y \not\perp x_S$) that is (directly or indirectly) exploited by the model


- features x_S not being causal for the prediction $\nRightarrow v(S) = 0$
 - e.g., if x_1 and x_2 are perfectly correlated, even if only x_1 has a nonzero coefficient, both are considered equally important
- under model optimality, links to mutual information or the conditional variance exist



SAGE - MARGINAL AND COND. SAMPLING

Example:

- $y = x_3 + \epsilon_y$
 $x_1 = \epsilon_1$
 $x_2 = x_1 + \epsilon_2$
 $x_3 = x_2 + \epsilon_3$ (all ϵ_j i.i.d.)
- Causal DAG:
 $x_1 \rightarrow x_2 \rightarrow x_3 \rightarrow y$
- Fitted LM:
 $\hat{f} \approx 0.95x_3 + 0.05x_2$



SAGE - MARGINAL AND COND. SAMPLING

Example:

- $y = x_3 + \epsilon_y$
 $x_1 = \epsilon_1$
 $x_2 = x_1 + \epsilon_2$
 $x_3 = x_2 + \epsilon_3$ (all ϵ_j i.i.d.)



- Causal DAG:

$$x_1 \rightarrow x_2 \rightarrow x_3 \rightarrow y$$

- Fitted LM:


$$\hat{f} \approx 0.95x_3 + 0.05x_2$$

- Marginal $v(j)$ are only nonzero for features that are used by $\hat{f}$
- Conditional $v(j)$ are also nonzero for features that are not used by $\hat{f}$ (e.g., due to correlation)
- For conditional value function v , the difference $v(-j \cup j) - v(-j)$ quantifies the unique contribution of x_j over remaining features x_{-j}
 $\Rightarrow$ Since $y \perp\!\!\!\perp x_1, x_2 | x_3$, only $v(\{1, 2, 3\}) - v(\{1, 2\})$ is nonzero (i.e., for feature $j = 3$)



SAGE VALUE FUNCTIONS VS. SAGE VALUES

SAGE value function $v(S)$: measure contribution of a specific feature set over the empty coalition




SAGE VALUE FUNCTIONS VS. SAGE VALUES

SAGE value function $v(S)$: measure contribution of a specific feature set over the empty coalition

SAGE values ϕ_j : fair attribution of importance

- can be computed by averaging the contribution of x_j over all feat orderings
- for feature permutation τ , the contribution of j in the set S_j^τ is given as $v(S_j^\tau \cup \{j\}) - v(S_j^\tau)$

Note: S_j^τ is the set of features preceding j in permutation τ



SAGE VALUE FUNCTIONS VS. SAGE VALUES


SAGE value function $v(S)$: measure contribution of a specific feature set over the empty coalition

SAGE values ϕ_j : fair attribution of importance

- can be computed by averaging the contribution of x_j over all feat orderings
- for feature permutation τ , the contribution of j in the set S_j^τ is given as $v(S_j^\tau \cup \{j\}) - v(S_j^\tau)$
Note: S_j^τ is the set of features preceding j in permutation τ



SAGE value approximation: Average over the contributions for M randomly sampled permutations

$$\phi_j = \frac{1}{M} \sum_{m=1}^M v(S_j^{\tau_m} \cup \{j\}) - v(S_j^{\tau_m})$$



INTERACTION EXAMPLE REVISITED


Recap: Data: $x_1, \dots, x_4$ uniformly sampled from $\{-1, 1\}$ and $y := x_1 x_2 + x_3 + \epsilon_Y$ with $\epsilon_Y \sim N(0, 1)$. Model: $\hat{f}(x) \approx x_1 x_2 + x_3$.



- PFI regards x_1, x_2 to be equally important as x_3
- Marginal SAGE fairly divides the contribution of the interaction x_1 and x_2

SAGE LOSS FUNCTIONS

When the loss-optimal model f^* is inspected using *conditional-sampling* based SAGE value functions, interesting links exist.




SAGE LOSS FUNCTIONS

When the loss-optimal model f^* is inspected using *conditional-sampling* based SAGE value functions, interesting links exist.

For cross-entropy loss:

- value function is the mutual information: $v_{f^*}(S) = I(y; x_S)$
- surplus contribution of a feature x_j is the conditional mutual information:
$$v_{f^*}(S \cup \{j\}) - v_{f^*}(S) = I(y, x_j | x_S)$$



SAGE LOSS FUNCTIONS


When the loss-optimal model f^* is inspected using *conditional-sampling* based SAGE value functions, interesting links exist.

For cross-entropy loss:

- value function is the mutual information: $v_{f^*}(S) = I(y; x_S)$
- surplus contribution of a feature x_j is the conditional mutual information:
$$v_{f^*}(S \cup \{j\}) - v_{f^*}(S) = I(y, x_j | x_S)$$

For MSE loss:

- value function is the expected reduction in variance given knowledge of the features x_S : $v_{f^*}(S) = \text{Var}(y) - \mathbb{E}[\text{Var}(y|x_S)]$
- surplus contribution is the respective reduction over x_S :
$$v_{f^*}(S \cup \{j\}) - v_{f^*}(S) = \mathbb{E}[\text{Var}(y|x_S)] - \mathbb{E}[\text{Var}(y|x_{S \cup \{j\}})]$$




IMPLICATIONS MARGINAL SAGE VALUES

Can we gain insight into whether the ...


❶ feature x_j is causal for the prediction?

- for all coalitions S , $v(j \cup S) - v(S)$ can only be nonzero if $x_j \rightarrow \hat{f}(x)$ (as for PFI)
 $\rightsquigarrow \phi_j$ is only nonzero if x_j is causal for the prediction
- $v(j \cup S) - v(S)$ may be zero due to indep. $x_j \perp\!\!\!\perp y | x_S$ (as for PFI)
 $\rightsquigarrow \phi_j$ may be zero although the feature is causal for the prediction



IMPLICATIONS MARGINAL SAGE VALUES


Can we gain insight into whether the ...



- ② feature x_j contains prediction-relevant information about y
 - value functions may be nonzero despite independence due to extrapolation (as for PFI)
 - $\rightsquigarrow \phi_j$ may be nonzero without x_j being dependent with y
 - value functions may be zero despite x_j containing prediction-relevant information due to underfitting (as for PFI)
 - $\rightsquigarrow \phi_j$ may be zero although prediction-relevant information contained

IMPLICATIONS MARGINAL SAGE VALUES

Can we gain insight into whether the ...




- ③ model requires access to x_j to achieve its prediction performance?
 - like PFI, in general marginal value functions do not allow insight into unique contribution $\rightsquigarrow$ no insight from ϕ_j

IMPLICATIONS CONDITIONAL SAGE VALUES

Can we gain insight into whether the ...


❶ feature $\mathbf{x}_j$ is causal for the prediction?

- value funcs may be nonzero although feature is not directly used by $\hat{f}$
 $\leadsto$ nonzero ϕ_j does not imply $\mathbf{x}_j \rightarrow \hat{y}$
- value functions may be zero although feature may be used by the model,
e.g. if feature is independent with y and all other features
 $\leadsto$ zero ϕ_j does not imply $\mathbf{x}_j \not\rightarrow \hat{y}$



IMPLICATIONS CONDITIONAL SAGE VALUES

Can we gain insight into whether the ...




② feature $\mathbf{x}_j$ contains prediction-relevant information about y ?

- e.g. for cross-entropy optimal $\hat{f}$, $v(j)$ measures mutual info. $I(y; x_j) \rightsquigarrow$ prediction-relevance implies nonzero ϕ_j
- $x_j \perp\!\!\!\perp y$ does not imply $x_j \perp\!\!\!\perp y|x_S$ and consequently does not imply $v(j \cup S) - v(S) = 0 \rightsquigarrow \phi_j$ may be nonzero although $\mathbf{x}_j \perp\!\!\!\perp y$

IMPLICATIONS CONDITIONAL SAGE VALUES


Can we gain insight into whether the ...



- ③ model requires access to x_j to achieve its prediction performance?
 - e.g. for cross-entropy optimal $\hat{f}$, surplus contrib. $v(j \cup -j) - v(-j)$ captures the conditional mutual information $I(y; x_j | x_{-j})$
 - $\leadsto \phi_j$ is nonzero for features with unique contribution
 - $x_j \perp\!\!\!\perp y | x_{-j}$ does not imply $x_j \perp\!\!\!\perp y | x_S$ (cond. w.r.t. to arbitrary coalitions S)
 - $\leadsto \phi_j$ may be nonzero although feature has no unique contrib.

DEEP DIVE: SHAPLEY AXIOMS FOR SAGE


The Shapley axioms can be translated into properties of SAGE. The interpretation depends on whether conditional or marginal sampling is used.



Shapley property $\implies$	conditional SAGE property
efficiency	$\sum_{i=1}^p \phi_j(\nu) = \mathcal{R}(\hat{f}_\emptyset) - \mathcal{R}(\hat{f})$
symmetry	$x_j = x_i \implies \phi_i = \phi_j$
linearity	ϕ_j expectation of per-instance conditional SHAP applied to model loss
monotonicity	given models f, f' , if $\forall S$: $v_f(S \cup j) - v_f(S) \geq v_{f'}(S \cup j) - v_{f'}(S)$ then $\phi_j(v_f) \geq \phi_j(v_{f'})$
dummy	if $\forall S : \hat{f}(x) \perp\!\!\!\perp x_j x_S \Rightarrow \phi_j = 0$

DEEP DIVE: SHAPLEY AXIOMS FOR SAGE

The Shapley axioms can be translated into properties of SAGE. The interpretation depends on whether conditional or marginal sampling is used.



Shapley property $\implies$	marginal SAGE property
efficiency	$\sum_{i=1}^p \phi_j(\nu) = \mathcal{R}(\hat{f}_\emptyset) - \mathcal{R}(\hat{f})$
symmetry	no intelligible implication
linearity	ϕ_j expectation of per-instance marginal SHAP applied to model loss
monotonicity	given models f, f' , if $\forall S :$ $v_f(S \cup j) - v_f(S) \geq v_{f'}(S \cup j) - v_{f'}(S)$ then $\phi_j(v_f) \geq \phi_j(v_{f'})$
dummy	model invariant to $x_j \Rightarrow \phi_j = 0$