

Interpretable Regional Descriptors


Susanne Dandl

30.04.2026

Counterfactual vs. Semi-factual Explanations


Counterfactual Explanations

Minimal changes in a few features required for changing a prediction



Semi-factual Explanations


Maximal changes in a few features required for **not** changing a prediction



Counterfactual vs. Semi-factual Explanations


Counterfactual Explanations

Minimal changes in a few features required for changing a prediction




Semi-factual Explanations

Maximal changes in a few features required for **not** changing a prediction



Motivation for Interpretable Regional Descriptors


How can the features be changed, such that the model still rates the credit as a moderate risk?



Feature	Customer	IRD	1-dim IRD
sex	female	{female, male}	{female, male}
saving.accounts	little	{little, moderate, rich}	{little, moderate, rich}
purpose	car	{car, radio/TV, furniture, others}	{car, radio/TV, furniture, others}
age	22	[19, 22]	[19, 75]
job	skilled	{skilled, highly skilled}	{unskilled, skilled, highly skilled}
housing	rent	{rent}	{own, free, rent}
checking.account	moderate	{little, moderate}	{little, moderate}
credit.amount	4000	[4000, 5389]	[2127, 8424]
duration	30	[26, 33]	[6, 44]

Motivation for Interpretable Regional Descriptors

How can the features be changed, such that the model still rates the credit as a moderate risk?




Feature	Customer	IRD	1-dim IRD
sex	female	{female, male}	{female, male}
saving.accounts	little	{little, moderate, rich}	{little, moderate, rich}
purpose	car	{car, radio/TV, furniture, others}	{car, radio/TV, furniture, others}
age	22	[19, 22]	[19, 75]
job	skilled	{skilled, highly skilled}	{unskilled, skilled, highly skilled}
housing	rent	{rent}	{own, free, rent}
checking.account	moderate	{little, moderate}	{little, moderate}
credit.amount	4000	[4000, 5389]	[2127, 8424]
duration	30	[26, 33]	[6, 44]

- 1 Provide a set of semi-factual explanations to justify a prediction (“even-if” statements)
“Even if you increase your savings and have a longer payback period, your credit would still be of moderate risk”

Motivation for Interpretable Regional Descriptors

How can the features be changed, such that the model still rates the credit as a moderate risk?




Feature	Customer	IRD	1-dim IRD
sex	female	{female, male}	{female, male}
saving.accounts	little	{little, moderate, rich}	{little, moderate, rich}
purpose	car	{car, radio/TV, furniture, others}	{car, radio/TV, furniture, others}
age	22	[19, 22]	[19, 75]
job	skilled	{skilled, highly skilled}	{unskilled, skilled, highly skilled}
housing	rent	{rent}	{own, free, rent}
checking.account	moderate	{little, moderate}	{little, moderate}
credit.amount	4000	[4000, 5389]	[2127, 8424]
duration	30	[26, 33]	[6, 44]

2 Help with model auditing

“If male customers are not part of the IRD this can be an indicator for pointwise biases”

Motivation for Interpretable Regional Descriptors

How can the features be changed, such that the model still rates the credit as a moderate risk?




Feature	Customer	IRD	1-dim IRD
sex	female	{female, male}	{female, male}
saving.accounts	little	{little, moderate, rich}	{little, moderate, rich}
purpose	car	{car, radio/TV, furniture, others}	{car, radio/TV, furniture, others}
age	22	[19, 22]	[19, 75]
job	skilled	{skilled, highly skilled}	{unskilled, skilled, highly skilled}
housing	rent	{rent}	{own, free, rent}
checking.account	moderate	{little, moderate}	{little, moderate}
credit.amount	4000	[4000, 5389]	[2127, 8424]
duration	30	[26, 33]	[6, 44]

3 Identify whether a feature affects a prediction locally

“Compared to credit amount or duration, savings or purpose seem less relevant because their IRD encompasses the entire observed feature ranges”

General Task




Given: ML model $\hat{f} : \mathcal{X} \rightarrow \mathbb{R}$, point of interest $\mathbf{x}'$

Aim: Identify the largest hyperbox B' among all hyperboxes B which cover $\mathbf{x}'$ and only contain observations with predictions in

$$Y' = [\hat{f}(\mathbf{x}') - \epsilon_L, \hat{f}(\mathbf{x}') + \epsilon_H], \epsilon_L, \epsilon_H \in \mathbb{R}_{\geq 0}$$

General Task



Given: ML model $\hat{f} : \mathcal{X} \rightarrow \mathbb{R}$, point of interest $\mathbf{x}'$

Aim: Identify the largest hyperbox B' among all hyperboxes B which cover $\mathbf{x}'$ and only contain observations with predictions in


$$Y' = [\hat{f}(\mathbf{x}') - \epsilon_L, \hat{f}(\mathbf{x}') + \epsilon_H], \epsilon_L, \epsilon_H \in \mathbb{R}_{\geq 0}$$

Optimization problem:

$$\begin{aligned} B' \in \arg \max_B & \underbrace{\mathbb{P}(\mathbf{x} \in B | \mathbf{x} \in \mathcal{X})}_{=: \text{coverage}(B)} \\ \text{s.t. } \underbrace{\mathbb{I}(\mathbf{x}' \in B)}_{=: \text{locality}(B)} &= 1 \text{ and } \underbrace{\mathbb{P}(\hat{f}(\mathbf{x}) \in Y' | \mathbf{x} \in B)}_{=: \text{precision}(B)} = 1 \end{aligned}$$

General Framework - 4 Steps

- 1 Definition of search space




2

3

4

General Framework - 4 Steps

- 1 Definition of search space




Theorem

For any B' solving
the above
optimization task:

$$B' \subseteq \overline{B}$$

General Framework - 4 Steps



1 Definition of search space



Theorem

For any B' solving
the above
optimization task:
 $B' \subseteq \bar{B}$

2 Selection of data set




3

4

General Framework - 4 Steps



1 Definition of search space



Theorem

For any B' solving
the above
optimization task:
 $B' \subseteq \bar{B}$

2 Selection of data set




3

4

RQ: What effect does sampled data have compared to training data?

General Framework - 4 Steps



1 Definition of search space





Theorem

For any B' solving
the above
optimization task:
 $B' \subseteq \bar{B}$

2 Selection of data set




3 Initialization of box



4

General Framework - 4 Steps



1 Definition of search space





Theorem

For any B' solving the above optimization task:
 $B' \subseteq \bar{B}$



2 Selection of data set



3 Initialization of box



4 Optimization of box



Review of Hyperbox Methods

Desiderata:

- Interpretable: return a single hyperbox
- Model-agnostic: work for regression/classification models and various feature types
- Sparse: adhere to optional user-defined immutability constraints

Review of Hyperbox Methods

	Objectives			Desiderata		
	Coverage	Precision	Locality	Interpretable	Agnostic	Sparse
Level set methods						
PBnB (Zabinsky et al. (2011, 2020))	✓	✓	×	×	✓	×
Data mining methods						
MaxBox (Eckstein (2002))	✓	✓	×	✓	×	×
PRIM (Friedman & Fisher (1999))	×	×	×	✓	×	×
Local interpretation methods						
Anchors (Ribeiro et al. (2018))	✓	✓	✓	✓	×	×
MAIRE (Sharma et al. (2021))	✓	✓	✓	✓	×	×
LORE (Guidotti et al. (2019, 2022))	×	×	✓	✓	✓	×
Interpretable classifiers						
Column generation (Dash et al. (2018))	✓	✓	×	×	✓	×

Review of Hyperbox Methods

	Objectives			Desiderata		
	Coverage	Precision	Locality	Interpretable	Agnostic	Sparse
Data mining methods						
MaxBox (Eckstein (2002))	✓	✓	×	✓	×	×
PRIM (Friedman & Fisher (1999))	×	×	×	✓	×	×
Local interpretation methods						
MAIRE (Sharma et al. (2021))	✓	✓	✓	✓	×	×

Baseline: Anchors (Ribeiro et al. (2018))

Eckstein et al. (2002). *The Maximum Box Problem and its Application to Data Analysis*. Computational Optimization and Applications.

Friedman, Fisher (1999). *Bump Hunting in High-dimensional Data*. Statistics and Computing.

Sharma et al. (2021). *MAIRE - A Model-agnostic Interpretable Rule Extraction Procedure for Explaining Classifiers*. Machine Learning and Knowledge Extraction.

Ribeiro et al. (2018). *Anchors: High-precision Model-agnostic Explanations*. AAAI.

Review of Hyperbox Methods

	Objectives			Desiderata		
	Coverage	Precision	Locality	Interpretable	Agnostic	Sparse
Data mining methods						
MaxBox (Eckstein (2002))	✓	✓	✓	✓	✓	✓
PRIM (Friedman & Fisher (1999))	✓	✓	✓	✓	✓	✓
Local interpretation methods						
MAIRE (Sharma et al. (2021))	✓	✓	✓	✓	✓	✓

Baseline: Anchors (Ribeiro et al. (2018))

Eckstein et al. (2002). *The Maximum Box Problem and its Application to Data Analysis*. Computational Optimization and Applications.

Friedman, Fisher (1999). *Bump Hunting in High-dimensional Data*. Statistics and Computing.

Sharma et al. (2021). *MAIRE - A Model-agnostic Interpretable Rule Extraction Procedure for Explaining Classifiers*. Machine Learning and Knowledge Extraction.

Ribeiro et al. (2018). *Anchors: High-precision Model-agnostic Explanations*. AAAI.

Review of Hyperbox Methods


	Objectives			Desiderata		
	Coverage	Precision	Locality	Interpretable	Agnostic	Sparse
Data mining methods						
MaxBox (Eckstein (2002))	✓	✓	✓	✓	✓	✓
PRIM (Friedman & Fisher (1999))	✓	✓	✓	✓	✓	✓
Local interpretation methods						
MAIRE (Sharma et al. (2021))	✓	✓	✓	✓	✓	✓

Baseline: Anchors (Ribeiro et al. (2018))

RQ: How do MaxBox, PRIM and MAIRE perform against each other?


Review of Hyperbox Methods

Input: initial box & data set




Post-Processing

Input: generated box & data sampler



Post-Processing

Input: generated box & data sampler



RQ: What effect does post-processing have?

Research Questions

- ① How do MaxBox, PRIM and MAIRE perform against each other?
- ② What effect does (double-in-size) sampled data have compared to training data as the underlying dataset?
- ③ What effect does post-processing have?

Quality Measures

- **Locality:** Does the box B' contain the data point $\mathbf{x}'$ we are interested in? $locality(B)$
- **Coverage:** How much of the feature space is covered by B' ? $\widehat{coverage}(B)$
- **Precision:** Do all data points $\in B'$ have the same prediction as $\mathbf{x}'$? $\widehat{precision}(B)$
- **Efficiency:** How often is $\hat{f}$ called by a method?

Quality Measures

- **Locality:** Does the box B' contain the data point $\mathbf{x}'$ we are interested in? $locality(B)$
- **Coverage:** How much of the feature space is covered by B' ? $\widehat{coverage}(B)$
- **Precision:** Do all data points $\in B'$ have the same prediction as $\mathbf{x}'$? $\widehat{precision}(B)$
- **Efficiency:** How often is $\hat{f}$ called by a method?
- **Maximality:** Is the box maximal, so that no more boxes can be added under full precision?
- **Robustness:** Does an IRD method ran multiple times on the same data point and model, will produce B' s that overlap?


Setup

- 6 datasets


Name	ID	Type	n	Cont.	Categ.
diabetes	37	binary	768	8	0
tic_tac_toe	50	binary	958	0	9
cmc	23	three-class	1473	2	7
vehicle	54	four-class	846	18	0
no2	886	regression	500	7	0
plasma_retinol	511	regression	315	10	3

- Models: hyperbox, logistic/multinomial/linear, neural network, random forest
- Methods: MaxBox, PRIM, MAIRE, Anchors (baseline), either with or without post-processing
- Underlying data: training data or double-in-size newly sampled data points within $\bar{B}$

Results




Results - RQ 1) Methods




RQ 1) All fulfill locality, PRIM highest coverage, MaxBox highest precision, equal no. calls to $\hat{f}$

Results - RQ 2) Underlying Data Set



RQ 2) Double-in-size sampled data higher coverage, precision w.r.t. sampled data

Results - RQ 3) Post-processing



RQ 3) Post-processed boxes higher coverage and precision, but also higher no. of calls to $\hat{f}$

Advantages & Disadvantages

Advantages:

- New form of explanation for single predictions
- Unifying framework that allows embedding existing methods
- Post-processing and sampling strategies to enhance quality of explanations

Advantages & Disadvantages

Advantages:

- New form of explanation for single predictions
- Unifying framework that allows embedding existing methods
- Post-processing and sampling strategies to enhance quality of explanations

Disadvantages:

- Explanation of model, not of DGP (ignores feature dependencies)
- Some methods are computationally demanding
- Restricted to tabular data